

```
In [1]: # KNN

import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
import seaborn as sns
```

```
In [2]: # Read in the data
raw_data = pd.read_csv('winequality-red.csv', sep=';')
```

```
In [3]: # raw_data.info()
# remove repeated data
raw_data = raw_data.drop_duplicates()
raw_data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 1359 entries, 0 to 1598
Data columns (total 12 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   fixed acidity                          1359 non-null   float64
1   volatile acidity                       1359 non-null   float64
2   citric acid                            1359 non-null   float64
3   residual sugar                         1359 non-null   float64
4   chlorides                             1359 non-null   float64
5   free sulfur dioxide                   1359 non-null   float64
6   total sulfur dioxide                   1359 non-null   float64
7   density                               1359 non-null   float64
8   pH                                     1359 non-null   float64
9   sulphates                             1359 non-null   float64
10  alcohol                               1359 non-null   float64
11  quality                               1359 non-null   int64
dtypes: float64(11), int64(1)
memory usage: 138.0 KB
```

```
In [4]: raw_data.describe()
```

```
Out[4]:
```

	fixed acidity	volatile acidity	citric acid	residual sugar	chlorides	free sulfur dioxide	total sulfur dioxide
count	1359.000000	1359.000000	1359.000000	1359.000000	1359.000000	1359.000000	1359.000000
mean	8.310596	0.529478	0.272333	2.523400	0.088124	15.893304	14.975048
std	1.736990	0.183031	0.195537	1.352314	0.049377	10.447270	9.673151
min	4.600000	0.120000	0.000000	0.900000	0.012000	1.000000	0.000000
25%	7.100000	0.390000	0.090000	1.900000	0.070000	7.000000	6.000000
50%	7.900000	0.520000	0.260000	2.200000	0.079000	14.000000	13.000000
75%	9.200000	0.640000	0.430000	2.600000	0.091000	21.000000	19.000000
max	15.900000	1.580000	1.000000	15.500000	0.611000	72.000000	28.000000

```
In [5]: X_train, X_test, y_train, y_test = train_test_split(raw_data.iloc[:, :-1], r
```

```
In [6]: # KNN
from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import accuracy_score


error_rate = []
for i in range(1, 40):

    knn = KNeighborsClassifier(n_neighbors = i)
    knn.fit(X_train, y_train)
    pred_i = knn.predict(X_test)
    error_rate.append(np.mean(pred_i != y_test))

plt.figure(figsize =(10, 6))
plt.plot(range(1, 40), error_rate, color ='blue',
         linestyle ='dashed', marker ='o',
         markerfacecolor ='red', markersize = 10)

plt.title('Error Rate vs. K Value')
plt.xlabel('K')
plt.ylabel('Error Rate')
```

```
Out[6]: Text(0, 0.5, 'Error Rate')
```



```
In [ ]:
```

```
In [7]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
import seaborn as sns
from sklearn.neighbors import KNeighborsClassifier
from sklearn.preprocessing import StandardScaler
```

```
In [8]: # read the data
raw_data = pd.read_csv('winequality-red.csv', sep=';')
```

```
In [9]: # remove outliers
for col in raw_data.columns:
    if col != 'quality':
        iqr = raw_data[col].quantile(0.75) - raw_data[col].quantile(
            0.25)
        upper_bound = raw_data[col].quantile(0.75) + 2.5 * iqr
        lower_bound = raw_data[col].quantile(0.25) - 2.5 * iqr
        raw_data = raw_data[(raw_data[col] > lower_bound) & (raw_data[col] < upper_bound)]

# save the cleaned data
raw_data.to_csv('cleaned_data.csv', index=False)
```

```
In [10]: # raw_data.info()
# remove repeated data
raw_data = raw_data.drop_duplicates()
raw_data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 1182 entries, 0 to 1598
Data columns (total 12 columns):
 #   Column              Non-Null Count  Dtype
---  -
 0   fixed acidity       1182 non-null   float64
 1   volatile acidity    1182 non-null   float64
 2   citric acid         1182 non-null   float64
 3   residual sugar      1182 non-null   float64
 4   chlorides           1182 non-null   float64
 5   free sulfur dioxide 1182 non-null   float64
 6   total sulfur dioxide 1182 non-null   float64
 7   density             1182 non-null   float64
 8   pH                 1182 non-null   float64
 9   sulphates          1182 non-null   float64
10   alcohol             1182 non-null   float64
11   quality             1182 non-null   int64
dtypes: float64(11), int64(1)
memory usage: 120.0 KB
```

```
In [11]: raw_data.describe()
```

Out [11]:

	fixed acidity	volatile acidity	citric acid	residual sugar	chlorides	free sulfur dioxide	total sulfur dioxide
count	1182.000000	1182.000000	1182.000000	1182.000000	1182.000000	1182.000000	1182.000000
mean	8.268613	0.525816	0.260398	2.255118	0.078864	15.692893	45.016941
std	1.674834	0.176989	0.189768	0.545431	0.015813	9.689748	30.281045
min	4.700000	0.120000	0.000000	0.900000	0.038000	1.000000	6.000000
25%	7.100000	0.390000	0.090000	1.900000	0.069000	8.000000	22.000000
50%	7.900000	0.520000	0.245000	2.200000	0.078000	14.000000	37.000000
75%	9.200000	0.640000	0.410000	2.500000	0.088000	21.000000	60.000000
max	14.300000	1.240000	0.760000	4.300000	0.136000	53.000000	152.000000

```
In [12]: # Standardization
scaler = StandardScaler()
raw_data.iloc[:, :-1] = scaler.fit_transform(raw_data.iloc[:, :-1])
```

```
In [13]: # train test split
from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(raw_data.iloc[:, :-1], r
```

```
In [14]: # # grid search on K and distance function
# from sklearn.model_selection import GridSearchCV
# k_range = list(range(1, 50))
# weight_options = ['uniform', 'distance']
# distance_metric = ['euclidean', 'manhattan', 'minkowski', 'cosine']
# param_grid = dict(n_neighbors=k_range, weights=weight_options, metric=dist
# knn = KNeighborsClassifier()
# grid = GridSearchCV(knn, param_grid, cv=10, scoring='accuracy', return_tra
# grid.fit(X_train, y_train)
#
# # print the best parameters
# print(grid.best_params_)
# print(grid.best_score_)
```

```

In [15]: train_score = []
test_score = []
for i in range(1, 40):

    knn = KNeighborsClassifier(n_neighbors = i)
    knn.fit(X_train, y_train)
    pred_i = knn.predict(X_test)
    test_score.append(knn.score(X_test, y_test))
    train_score.append(knn.score(X_train, y_train))

plt.figure(figsize =(10, 8))

plt.plot(range(1, 40), test_score, color ='blue', linestyle ='dashed', marker='o')


plt.plot(range(1, 40), train_score, color ='red', linestyle ='dashed', marker='o')

plt.xlabel('K')
plt.legend(['test score', 'train score'], loc ='upper right')

print("Maximum test score:",max(test_score),"at K =",test_score.index(max(test_score)))

```

Maximum test score: 0.6197183098591549 at K = 17



```
In [16]: # k = 17
classifier = KNeighborsClassifier(n_neighbors = 17)
classifier.fit(X_train,y_train)
```

```
Out[16]: ▼      KNeighborsClassifier
KNeighborsClassifier(n_neighbors=17)
```

```
In [17]: #Predicting the ouput from input data (x_train) and (y_train)
y_pred1 = classifier.predict(X_train)
y_pred2 = classifier.predict(X_test)
```

```
In [18]: from sklearn.metrics import accuracy_score, mean_squared_error


print("train score",accuracy_score(y_train, y_pred1))
print("test score",accuracy_score(y_test, y_pred2))
print("MSE",mean_squared_error(y_test, y_pred2))

y_test = np.array(y_test)

train score 0.6287787182587666
test score 0.6197183098591549
MSE 0.49859154929577465
```

```
In [19]: # visualization
correct = 0
one_bit_error = 0
two_bit_error = 0
threemore_bit_error = 0
print('Shap of y_pred: ', y_test.shape)
for i in range(len(y_pred2)):
    if y_pred2[i] == y_test[i]:
        correct += 1
    elif abs(y_pred2[i] - y_test[i]) == 1:
        one_bit_error += 1
    elif abs(y_pred2[i] - y_test[i]) == 2:
        two_bit_error += 1
    else:
        threemore_bit_error += 1
plt.figure(figsize=(12, 12))
plt.bar(['correct', '1 bit error', '2 bit error', '3 bit error'], [correct,
plt.title('Prediction of test data')
plt.xlabel('Prediction')
plt.ylabel('Number of data')
plt.show()

Shap of y_pred:  (355,)
```



In []: